

Microsoft : Classifying Cybersecurity Incidents with Machine Learning

Objective

1

Preprocess and Clean Data

Preprocess and clean large-scale incident data (1.3M rows).

2

Feature Engineering

Engineer features for better model performance.

3

Train and Evaluate Model

Train and evaluate a high-performing classification model.

4

Provide Insights

Provide interpretability and actionable insights from predictions.

Dataset Description

Key Details

Size: 1,297,443 rows × 39 columns

Key Features: Category,
IncidentGrade, EntityType, Hour,
DayOfWeek, etc.

Target Variable Distribution

BenignPositive: 2,054,774

TruePositive: 1,662,087

FalsePositive: 1,015,782

Missing Values Summary

Columns dropped (missing >50%):
ActionGrouped, ResourceType, etc.

Imputed numerical and categorical
columns with median/mode.

Preprocessing Steps

Data Cleaning

Removed duplicates (0 rows).

Handled missing values by imputation.

Outlier Removal

Used IQR method for numerical features.

Feature Engineering

Extracted temporal features: Year, Month, Hour, DayOfWeek.

Encoded categorical features with Label Encoding and One-Hot Encoding.

Scaling

Applied Min-Max Scaling to numerical features.

Choosing the Right Model

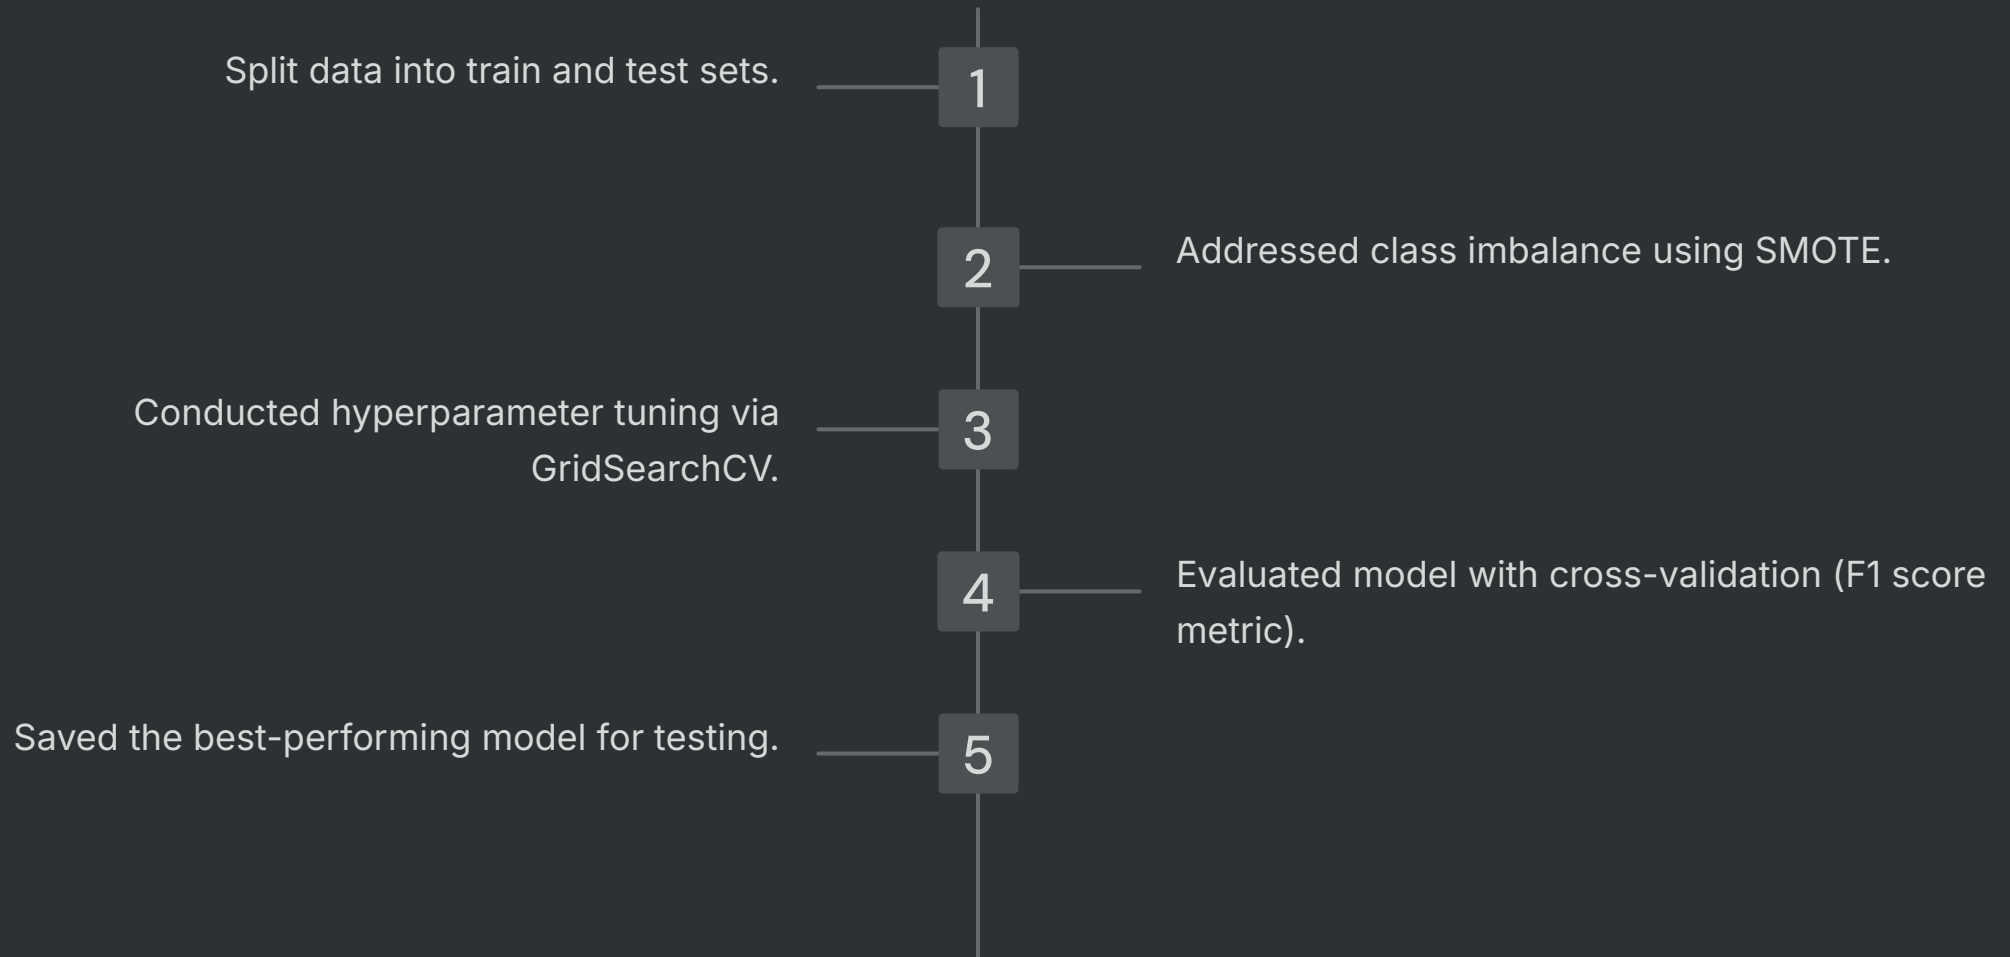
Model Chosen

After careful evaluation, we selected the **Random Forest Classifier** as our machine learning model.

Why Random Forest?

- Handles large, complex datasets efficiently
- Robust to overfitting through ensemble learning
- Provides valuable feature importance insights
- Performs exceptionally well on imbalanced data with SMOTE

Training Process



Test Dataset Workflow

1

Data Cleaning and Preprocessing (same as training).

2

Feature alignment with training dataset.

3

Loaded saved model for predictions.

4

Evaluated test performance.

Challenges Faced



Class Imbalance

Solved using SMOTE.



High Missing Values

Dropped columns (>50%) and imputed remaining.



Overfitting

Addressed with cross-validation and hyperparameter tuning.



Temporal Data Handling

Extracted features like Hour, DayOfWeek, etc.

Final Results

96%

Train F1 Score

59%

Test F1 Score

Reason for Difference:

Despite performing **cross-validation and hyperparameter tuning** with RandomSearch to prevent overfitting, the **model consistently overfit**. Running adjustments **took 3–4 hours due to the large dataset**, making local execution challenging. Applying **SMOTE further increased runtime** but did not resolve the overfitting issue.

Conclusion and Future Work

